



What Counts as Responsible Research in the Age of Generative AI? Norms, Limits, and Governance

Karl de Fine Licht

Associate Professor, Ethics and Technology
Chalmers University of Technology

Top Takeaways

1. AI capability is not plateauing. It is accelerating and reaching more people than ever.

Industry produced over 90% of notable frontier models in 2025, and several of those models now meet or exceed human baselines on PhD-level science questions, multimodal reasoning, and competition mathematics. On a key coding benchmark—SWE-bench Verified—performance rose from 60% to near 100% in a single year. Organizational adoption reached 88%, and 4 in 5 university students now use generative AI.

See: [Chapter 2: Technical Performance](#)

Stanford Human-Centered Artificial Intelligence (HAI)

The 2026 AI Index Report



Top Takeaways

1. AI capability is not plateauing. It is accelerating and reaching more people than ever.

Industry produced over 90% of notable frontier models in 2025, and several of those models now meet or exceed human baselines on PhD-level science questions, multimodal reasoning, and competition mathematics. On a key coding benchmark—SWE-bench Verified—performance rose from 60% to near 100% in a single year. Organizational adoption reached 88%, and 4 in 5 university students now use generative AI.


See: [Chapter 2: Technical Performance](#)

9. AI agents advanced from answering questions to completing tasks in 2025, though they still fail roughly one in three attempts on structured benchmarks.

On OSWorld, which tests agents on computer tasks across operating systems, accuracy rose from roughly 12% to 66.3%, within 6 percentage points of human performance.

OSWorld: accuracy

Source: Epoch AI, 2026 | Chart: 2026 AI Index report



Top Takeaways

1. AI capability is not plateauing. It is accelerating and reaching more people than ever.

Industry produced over 90% of notable frontier models in 2025, and several of those models now meet or exceed human baselines on PhD-level science questions, multimodal reasoning, and competition mathematics. On a key coding benchmark—SWE-bench Verified—performance rose from 60% to near 100% in a single year. Organizational adoption reached 88%, and 4 in 5 university students now use generative AI.

See: [Chapter 2: Technical Performance](#)

Open-world evaluations for measuring frontier AI capabilities

Sayash Kapoor^{1,*} Peter Kirgis¹ Andrew Schwartz^{1,2} Stephan Rahauer¹
J.J. Allaire¹ Rohi Bonmann¹ Magda Dubois¹ Gillian Hatfield¹
Andy Hall¹ Sara Hooker¹ Seth Lazar^{1,3} Steve Newman¹
Dimitris Papailiopoulos^{10,11} Shoshannah Tekofsky^{1,2} Helen Tuner^{1,2} Corinna Uthoff¹
Arvind Narayanan¹

¹Princeton University ²Centroware Labs ³Stanford University
⁴UK AI Security Institute ⁵Johns Hopkins University ⁶Adaption Labs
⁷Australian National University ⁸Golden Gate Institute for AI ⁹UW Madison
¹⁰Microsoft Research ¹¹AI Digest ¹²Georgetown University (CSET)

Abstract

Benchmark-based evaluation remains important for tracking frontier AI progress. But we argue that it can both overstate and understate real-world capability because it privileges tasks that are precisely specified, automatically graded, relatively easy to optimize for, and run with low budgets and short time horizons. We advocate for a complementary class of evaluations, which we term *open-world evaluations*: long-horizon, messy, real-world tasks assessed through small-sample qualitative analysis rather than benchmark-scale automation. We survey recent open-world evaluations, identify their strengths and limitations, and introduce CRUX (Collaborative Research for Updating AI Expectations), a project for conducting such evaluations on a regular basis. As a first instance, we task an AI agent with developing and publishing a simple iOS application to the Apple App Store. The agent completed the task with a single unnecessary manual intervention, suggesting that open-world evaluations can provide early warning of capabilities that may soon become widespread. We conclude with recommendations for designing and reporting open-world evaluations.

1 Introduction

Tracking and predicting the capabilities of frontier AI systems is an open methodological problem, and one whose stakes rise as these systems are deployed in increasingly consequential settings. The predominant approach is benchmarking, in which agents are scored on large suites of automatically graded tasks. Benchmark-based evaluations underpin much of the public discussion about AI progress. For example, METE's time horizon graph [1] has been widely cited by policy analysts, industry leaders, and safety-focused organizations as evidence of rapid capability growth, and has been used to argue both for accelerated deployment and for tighter oversight. Decisions about funding, regulation, and safety investment are increasingly being made on the basis of such measurements. Benchmarks, however, can both overestimate and underestimate progress. Constructing a benchmark requires tasks that are precisely specified and automatically verifiable. But any task specified precisely enough to benchmark is also specified precisely enough to optimize for, which allows agents to excel on such tasks without necessarily acquiring the underlying capability. Test sets on prominent benchmarks can also leak into training data, either directly or through training-set instances that closely resemble the held-out evaluation tasks. Conversely, low benchmark accuracy may result

*Correspondence to: {kapoor, pkirgis, rahauer, arvind}@princeton.edu

What we do with GenAI

1. Idea generation
2. Research design
3. Data production/collection
4. Analysis
5. Writing

What we do with GenAI

1. Idea generation
2. Research design
3. Data production/collection
4. Analysis
5. Writing

Repairing the Crowd: Creator Communication and the Interpretation of Post-Campaign Failure in Reward Crowdfunding

Management-journal working paper draft with integrated theory, tables, figures, and robustness tests

Abstract

Reward crowdfunding converts uncertain creative and entrepreneurial projects into public stakeholder relationships. Although prior research explains why campaigns raise funds, less is known about how backers interpret projects after funding, when delays, changes, and non-delivery become salient. I theorize post-campaign creator communication as relational repair: a scalable mechanism through which creators preserve stakeholder trust after expectations are threatened. Using 36,229 individual backer-project evaluations matched to objective Kickstarter traces, I find that creator updates after the estimated delivery date are positively associated with backer-rated success, willingness to back the creator again, happiness with the outcome, and perceived communication quality. Backer comments after the estimated delivery date have the opposite association, consistent with stakeholder voice and contestation. Reward nonreceipt strongly depresses evaluations, but communication remains meaningfully associated with interpretation. The paper contributes to crowdfunding research, stakeholder theory, and trust repair theory by showing that post-campaign success is not only an operational outcome but also a managed social judgment.

Keywords: crowdfunding; stakeholder theory; trust repair; entrepreneurship; platform governance; communication

Introduction

Crowdfunding campaigns are usually judged at the moment the campaign clock stops. A project either reaches its funding goal or it does not; the platform records the outcome, observers celebrate or ignore the campaign, and researchers often explain why some projects raised money while others failed. Yet for backers, creators, and platforms, the most

What we do with GenAI

1. Idea generation
2. Research design
3. Data production/collection
4. Analysis
5. Writing

[Home](#) > [International Journal for Educational Integrity](#) > [Article](#)

Maintaining research integrity in the age of GenAI: an analysis of ethical challenges and recommendations to researchers

Research | [Open access](#) | Published: 02 July 2025
Volume 21, article number 18, (2025) [Cite this article](#)

● You have full access to this open access article

[Download PDF](#) ↓

[Save article](#)

[Sonja Bjelobaba](#) ✉, [Lorna Waddington](#), [Mike Perkins](#), [Tomáš Foltrýnek](#), [Sabuj Bhattacharyya](#) & [Deborah Weber-Wulff](#)

11k Accesses 15 Citations 35 Altmetric 2 Mentions [Explore all metrics](#) →

Abstract

Background


This paper is a practice-informed rapid review that maps the complex ethical challenges arising from the growing use of Generative Artificial Intelligence (GenAI) tools across the research life-cycle. Earlier research primarily focused on academic integrity concerns related to students' use of GenAI tools; however, limited information is available on the impact of GenAI on academic research. This study aims to examine the ethical concerns arising from the use of GenAI across different phases of research and explores potential strategies to encourage its ethical use for research purposes.

Methods

We selected one or more GenAI platforms applicable to various research phases (e.g. *developing research questions, conducting literature reviews, processing data, and*

How we react on GenAI use

- Lower quality
- Less original
- Spend less time on it




What should we say about this?

- Wonderful?
- Deeply problematic?
- Something in between?



How should we react?

- What are the alternatives?
 1. Prohibit use
 2. Allow some use, hold ourselves accountable (research)
 3. Allow some use, hold ourselves accountable (research, promotions, etc.), hold others accountable



What alternative should we choose?

- Outcomes + Values
- Outcomes: focus not only on GenAI

May 30, 2024 Product

Introducing ChatGPT Edu

An affordable offering for universities to responsibly bring AI to campus.



Values of researchers doing research:

1. Downstream: Societal benefits (“curing cancer”)
2. Epistemic: New knowledge (e.g., high-quality papers, robust findings)
3. Practice: Inquiry as a valuable human activity (e.g., curiosity, judgment, critical discussion)
4. Civic-pedagogical: Better teachers and public “engagers”

Maintaining values

1. Allow some use
2. Hold ourselves accountable
3. Revise our institutional setups
4. Hold others accountable



